



# TensorFlow and Keras

---

[statinfer.com](https://statinfer.com)

# Contents

- Deep Learning frameworks
- What is TensorFlow
- Key terms in TensorFlow
- Working with TensorFlow
- Regression Model building
- Keras Introduction
- Keras Advantages
- Working with Keras
- MNIST on Keras

# Limitations of Machine Learning tools

- Python, R and SAS work really well for solving the predictive modelling and machine learning problems
- The libraries like “sklearn” are sufficient for building regression models, trees, random forest and boosting models.
- But these tools have limited deep neural networks libraries
- What are the tools/frameworks for deep learning algorithms?

# Deep Learning Frameworks

- TensorFlow (by Google)
- Torch (by Facebook)
- Caffe (by UC Berkeley)
- Theano (Old version of TensorFlow)
- MxNet (by Amazon)
- CNTK (by Microsoft)
- Paddle (by Baidu)

theano



*dmlc*  
***mxnet***



# TensorFlow

- TensorFlow was developed by the Google Brain team for internal use.
- It was released under the Apache 2.0 open source license on November 9, 2015

DEC 1, 2015 @ 01:34 PM

4,801

The Little Black Book of I

## Reasons Why Google's Latest AI-TensorFlow is Open Sourced

CADE METZ BUSINESS 11.09.15 09:00 AM

**GOOGLE JUST OPEN SOURCED  
TENSORFLOW, ITS ARTIFICIAL  
INTELLIGENCE ENGINE**

statinfer.com

# TensorFlow

- Most popular among all Deep learning frameworks.
- TensorFlow works really well with matrix computations - All the deep learning algorithms are highly calculation intensive.
- Scalable to multi-CPU's and even GPU's
- Can handle almost all type of deep networks, be it ANN or CNN or RNNs



# Working with TensorFlow

- Has Python API and python is very easy install and to work on.
- We can use numpy to build all the models from scratch. But TensorFlow does it better by providing function to do it easily.
- TensorFlow has one of the best documentation and great community support as of now.

# TensorFlow Installation

```
!pip install tensorflow
```



# Some key terms


- Tensors
- Dimensions
- Computational graphs
- Nodes and Edges

# Tensor

A vector/array is a collection of elements (scalars)




A matrix is a two dimensional vector



Tensor means data

A Tensor is a multi-dimensional vector.



# Tensor

400
350
450
600
980
200

A vector indicating the product price(say smart phones)  
A tensor of dimension [6]

400	4
350	5
450	3
600	4
980	5
200	2

A matrix indicating the product price and rating  
A tensor of dimension [6,2]

400	4
350	5
450	3
600	4
980	5
200	2

A tensor indicating the product price and rating for last three years  
A tensor of dimension [6,2,3]


A tensor indicating the product price and rating for last three years in two countries  
A tensor of dimension [6,2,3,2]

400	4
350	5
450	3
600	4
980	5
200	2

# Tensor for images



- A colour image is represented as a three dimensional tensor
- [Width, Height, Colour]
- The colour component depth 3, Red, Green and Blue

16X16 pixels colour image  
A tensor with dimensions  
[16,16,3]



# Computational graphs

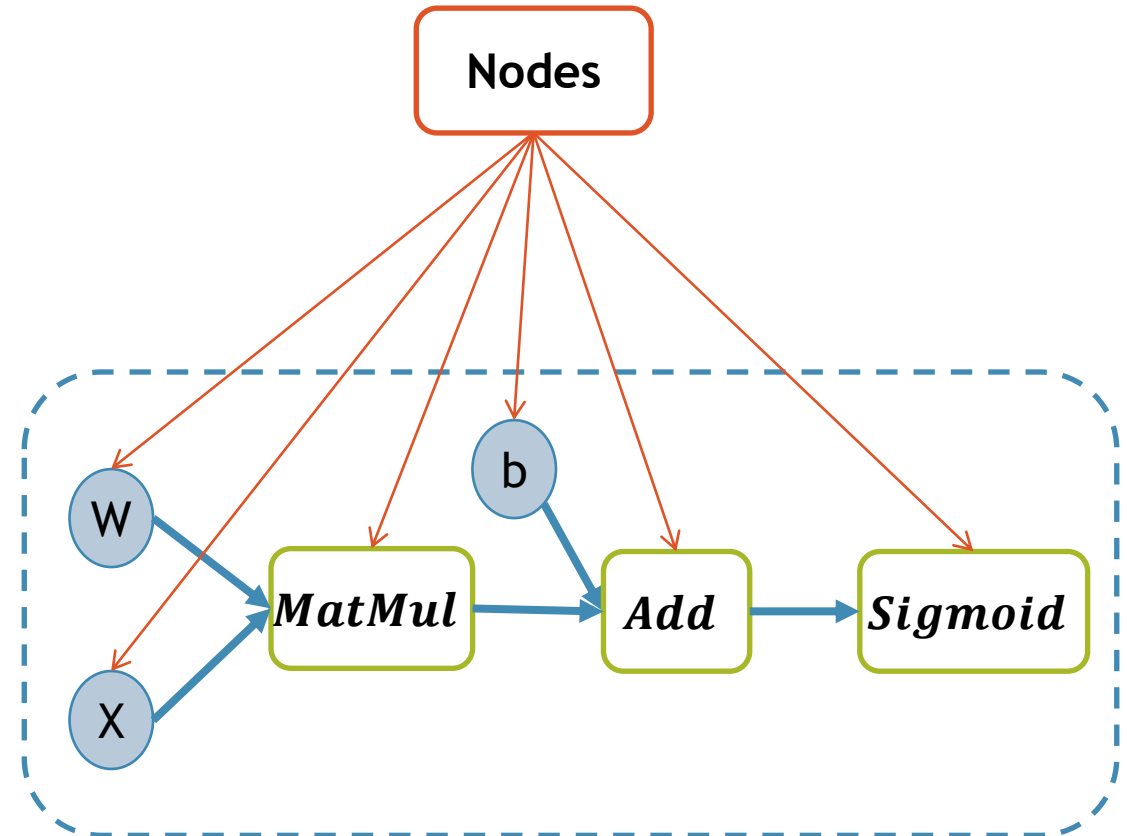
- Inside TensorFlow computations are represented using computational graphs
- You can call it as data flow graph. As sequence of operations on tensors(data)
- It has nodes and edges



# Computation Graphs-Nodes

- **Nodes**

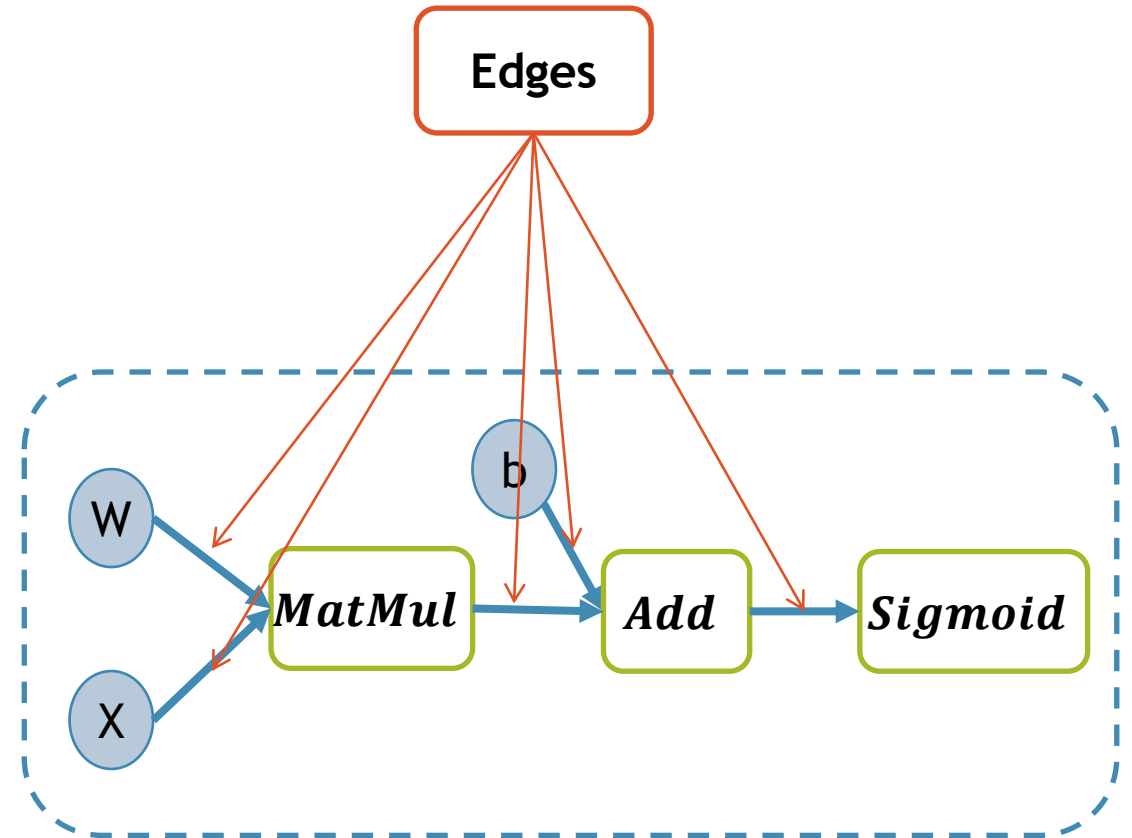
1. Data and Operations
2. Operations which have any number of inputs and outputs.
3. Variables/Tensors are also represented by nodes




# Computation Graphs- Edges

- **Edges :**

- Data flow direction
- Flow of tensors between nodes



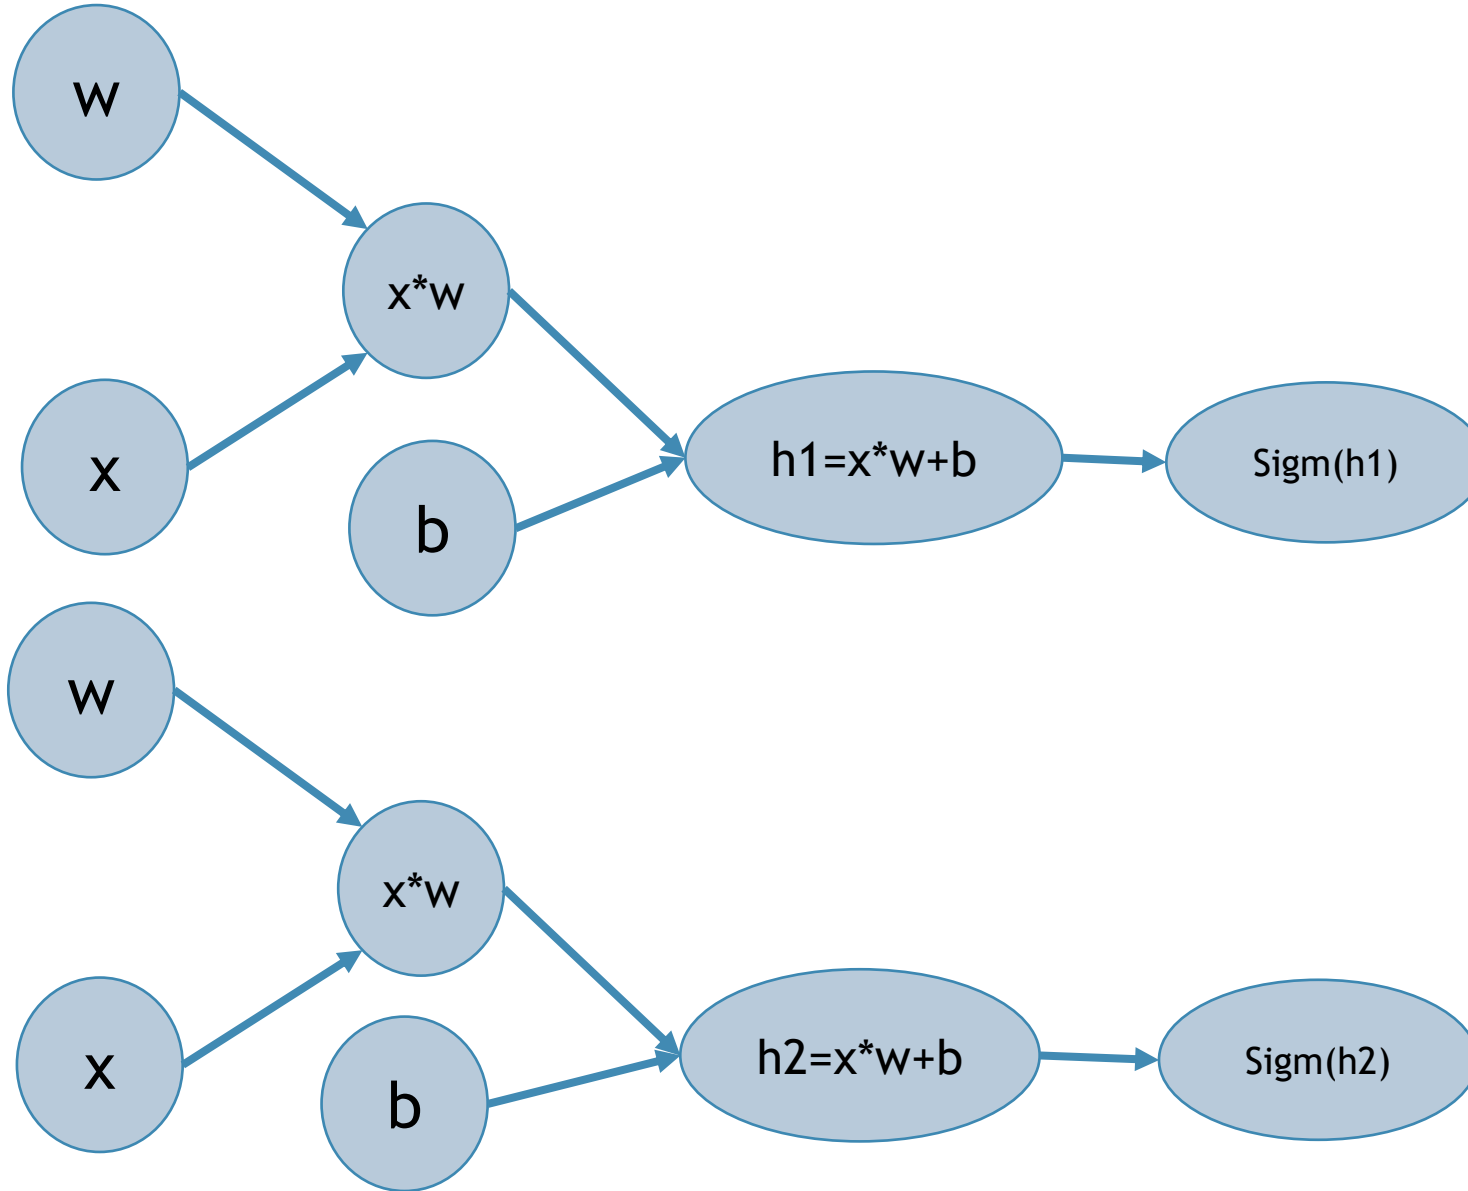
# Computation Graphs



- Computational graphs are particularly useful if the operations are complex
- TensorFlow computations define a computation graph that has no numerical value until it is evaluated!



# Computation Graphs



# Why Computation Graphs?

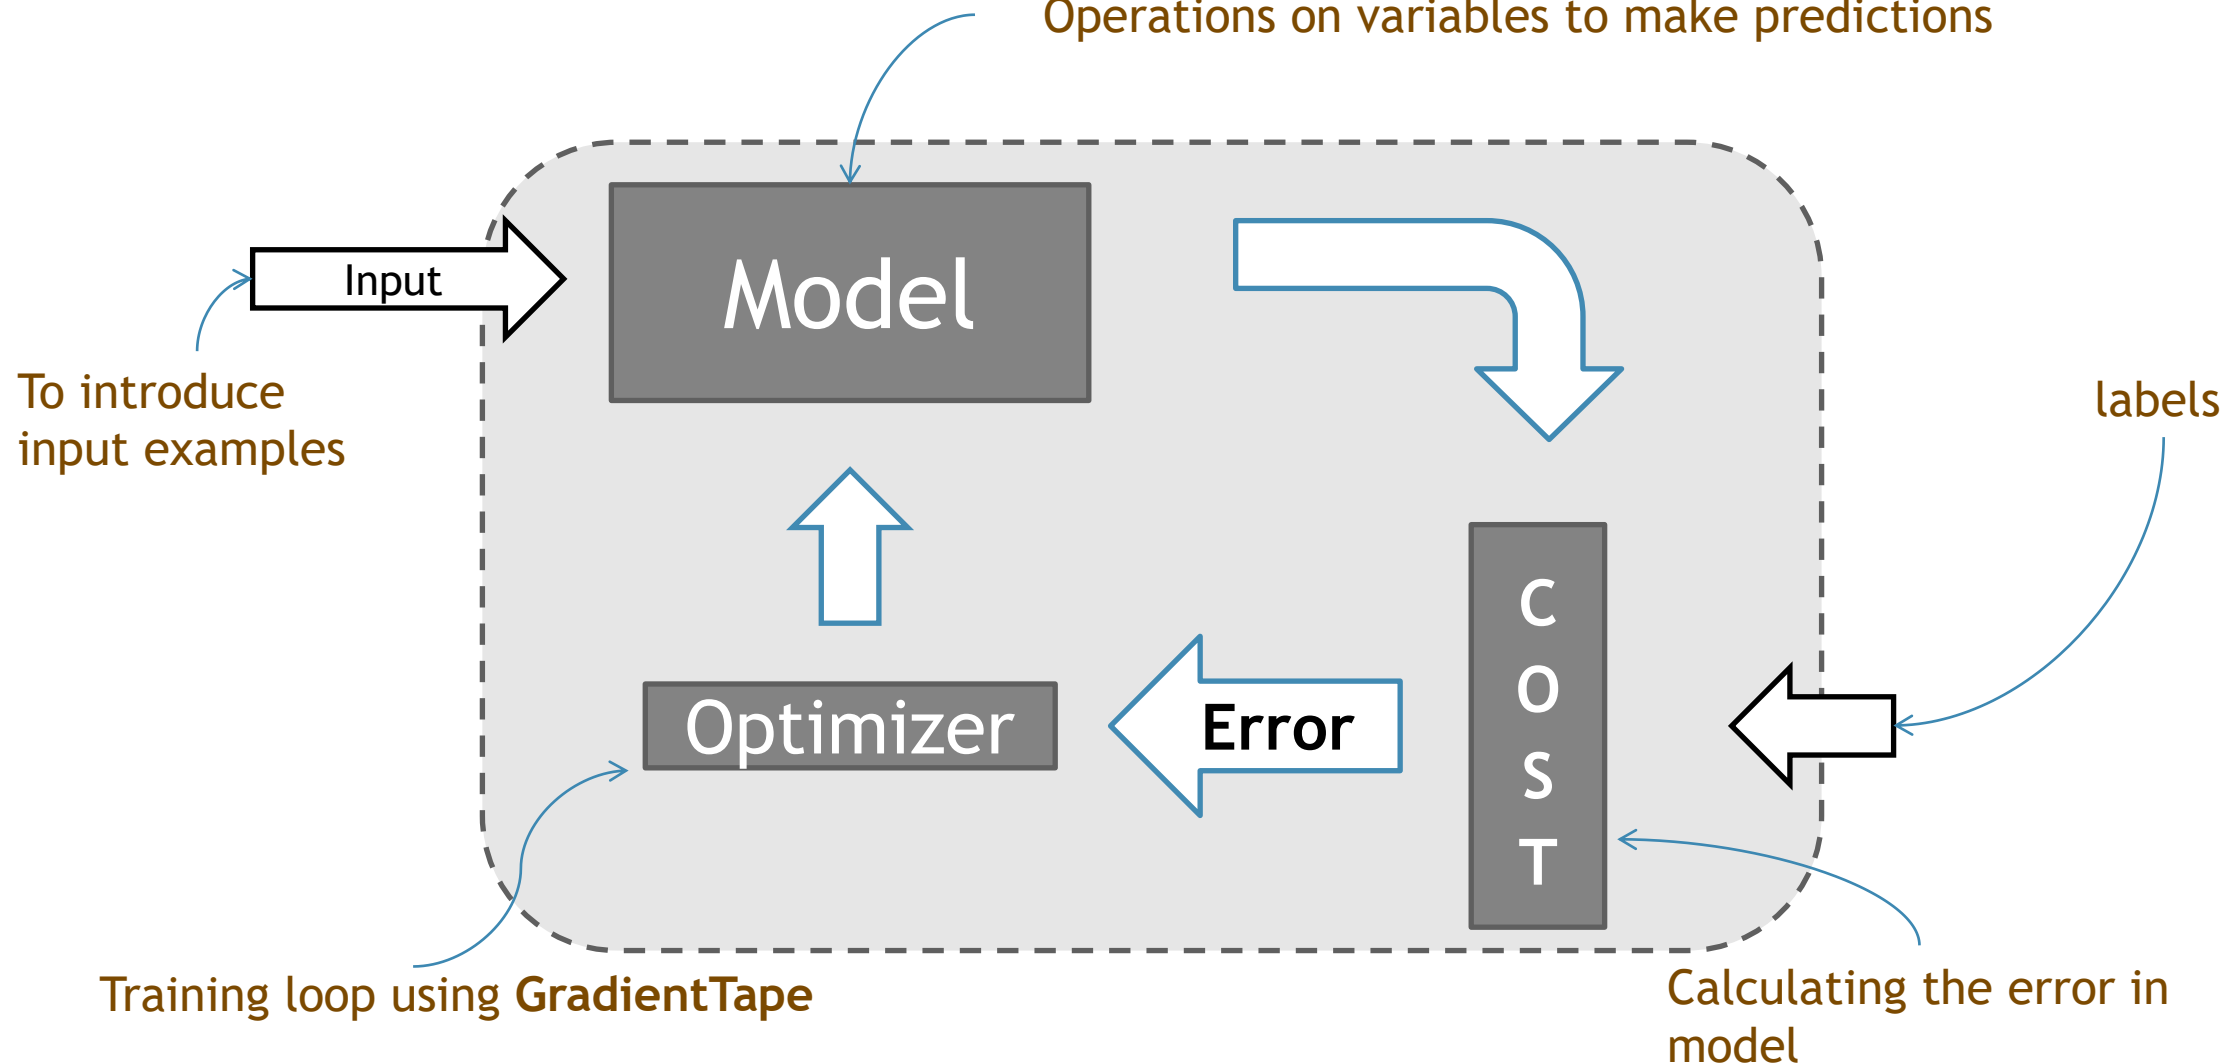
- Saves time by independently running the subgraphs that contribute to the final computation
- While training deep learning models, partial derivatives and chain rule applications are handled efficiently using these computational graphs.
- Break computation into small, differential pieces to facilitates auto-differentiation.
- Facilitate distributed computation, spread the work across multiple CPUs, GPUs, or devices.

# Model building on TensorFlow

- Write model equation  $Y = \text{sig}(X*W + b)$  or  $Y = X*W + b$
- Initialize the parameters (weights)
- Define loss function(cost)
- Write training loop using gradient tape
  - Optimise the loss function to find the best parameter estimates using the gradients

# Regression Model Training

Operations on variables to make predictions

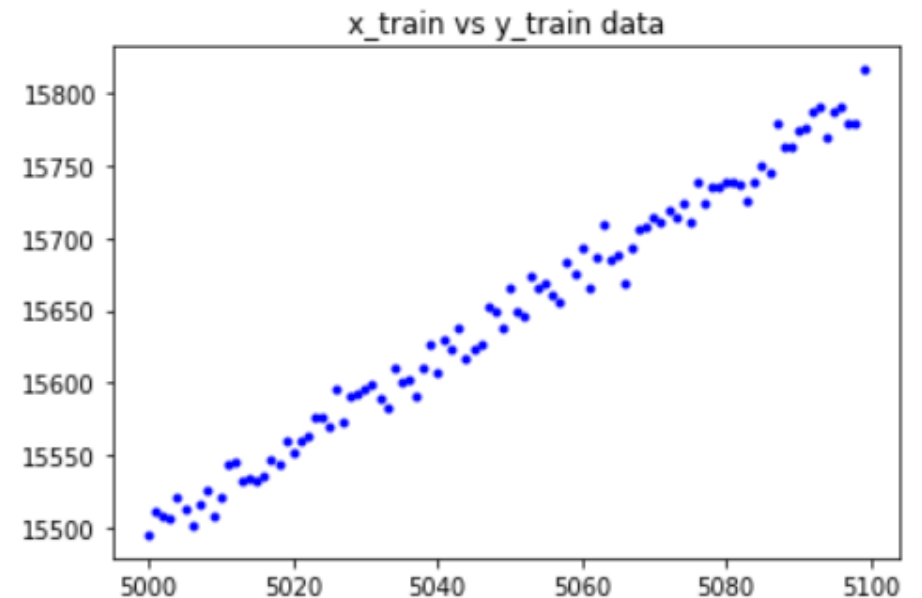


# LAB: Linear Regression on TensorFlow

- Use code to generate the data for  $x$  and  $y$ .
- Build a regression model in TensorFlow to find  $W$  and  $b$
- Verify the weight with the original data

# Code: Linear Regression on TensorFlow

```
#This step is for data creation, x and y  
import numpy as np  
x_train= np.array(range(5000,5100)).reshape(-1,1)  
  
y_train=[3*i+np.random.normal(500, 10) for i in x_train]  
  
import matplotlib.pyplot as plt  
plt.title("x_train vs y_train data")  
plt.plot(x_train, y_train, 'b.')  
plt.show()
```



# Code: Linear Regression on TensorFlow

```
#Model  $y=X*W + b$ 
#Model function
def output(x):
    return W*x + b

#Loss function Reduce mean square
def loss_function(y_pred, y_true):
    return tf.reduce_mean(tf.square(y_pred - y_true))

#Initialize Weights
W = tf.Variable(tf.random.uniform(shape=(1, 1)))
b = tf.Variable(tf.ones(shape=(1,)))

#Optimization
## Writing training/learning loop with GradientTape
learning_rate = 0.000000001
steps = 200 #epochs

for i in range(steps):
    with tf.GradientTape() as tape:
        predictions = output(x_train)
        loss = loss_function(predictions, y_train)
        dloss_dw, dloss_db = tape.gradient(loss, [W, b])
    W.assign_sub(learning_rate * dloss_dw)
    b.assign_sub(learning_rate * dloss_db)
    print(f"epoch : {i}, loss {loss.numpy()}, W : {W.numpy()}, b {b.numpy()}")
```

# LAB: Simple ANN Model

- Data:
- Our model function:

$$h = \text{Sigmoid}(Wx + b)$$

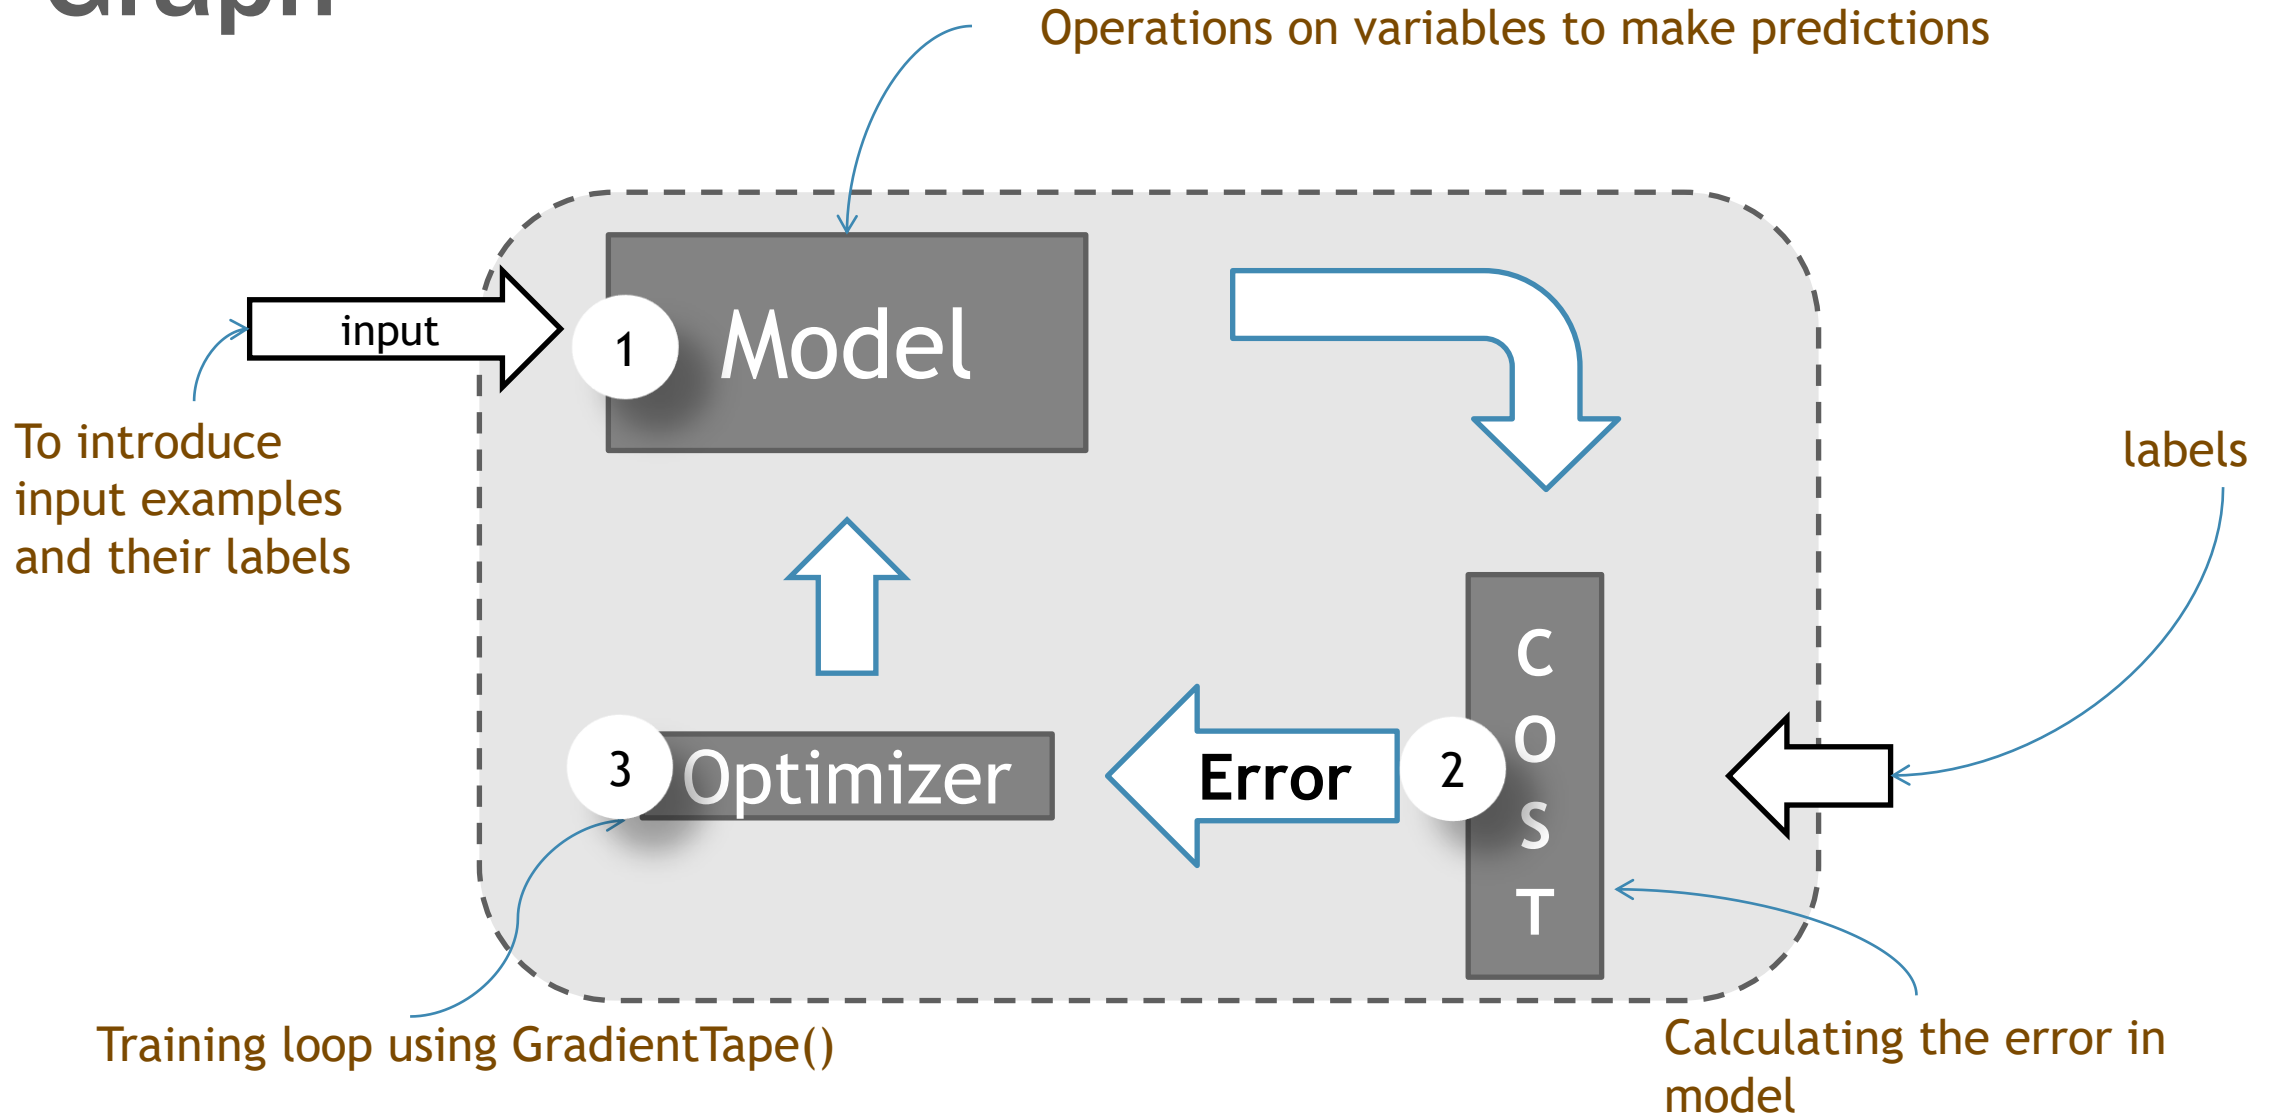
- The cost function:

$$J(W, b) = \frac{1}{N} \sum_{i=1}^N (y_i - h)^2$$

- Optimizer for back propagation



# Graph



# Lab: ANN Model– Single Perceptron


- Generate dummy x and y data
- Define our Model as Computation Graph
  - Place holders for X and y
  - W, b variables
  - Model output function
  - Cost function
  - Optimizer
- Run the graph as session, feeding data into placeholders

# TensorFlow Advantages



- complex deep learning computations easy
- Very fast compared to other frameworks
- Visualizations - Tensor Board
- GPU for faster computations

- Lot of low level coding, may take some time to get familiarity.



# Keras: TensorFlow made easy!!!

- Wrapper
  - Keras is a wrapper on top of TensorFlow.
  - High level API written in Python
- Easy
  - Less lines of code.
  - Easy to learn and implement deep learning models
- Best
  - Wide ranging options
  - Probably the best wrapper on top of TensorFlow



# Keras: TensorFlow made easy!!!

- Non-coders
  - Simple straight forward syntax
  - Provides detailed model summary statistics
  - Non-coders can start deep learning models with Keras
- Documentation
  - Good documentation on [keras.io](https://keras.io)
  - Good support from community and userbase



# Keras

## You have just found Keras.

Keras is a high-level neural networks API, written in Python and capable of running on top of **TensorFlow**, **CNTK**, or **Theano**. It was developed with a focus on enabling fast experimentation. *Being able to go from idea to result with the least possible delay is key to doing good research.*

Use Keras if you need a deep learning library that:

- Allows for easy and fast prototyping (through user friendliness, modularity, and extensibility).
- Supports both convolutional networks and recurrent networks, as well as combinations of the two.
- Runs seamlessly on CPU and GPU.

Read the documentation at **Keras.io**.

# Major steps in model building on Keras

1. Prepare your data.
2. Model Configuration
  - Add input layer to your model object
  - Add the hidden layers
  - Add the output layer
3. Compile the model object
4. Finally train the model

What are Layers?

# Sequence of Layers in the model

- Models building is done using sequence of layers
- The sequential model is a linear stack of layers.
- The first layer in the stack is “Input Layer” - Model receives the information on input shape
- The last layer is “Output Layer”. The model gets information on labels.
- We can add all the “model layers” in between. The model will prepare the weight parameters
- Lets see an example





# **LAB:** MNIST on Keras

---

Example of ANN on MNIST data using Tensorflow-Keras

# Code: MNIST on tf.Keras

- Importing our Keras from tensorflow

```
.....  
from tensorflow import keras  
from tensorflow.keras import layers
```

# Code: MNIST on Keras

- Keras's default MNIST data is in different format, make it a bit friendly.

```
# the data, shuffled and split between train and test sets
(X_train, Y_train), (X_test, Y_test) = keras.datasets.mnist.load_data()
num_classes=10
x_train = X_train.reshape(60000, 784)
x_test = X_test.reshape(10000, 784)
x_train = x_train.astype('float32')
x_test = x_test.astype('float32')
x_train /= 255
x_test /= 255
print(x_train.shape[0], 'train samples')
print(x_test.shape[0], 'test samples')
```

```
60000 train samples
10000 test samples
```

```
# convert class vectors to binary class matrices
y_train = keras.utils.to_categorical(Y_train, num_classes)
y_test = keras.utils.to_categorical(Y_test, num_classes)
```

Scaling the pixel  
values between 0 and  
1



Class vector needs to  
be converted into  
binary class matrices

# Code: MNIST on Keras

- Having a look at images using [matplotlib](#)

```
%matplotlib inline
import matplotlib.pyplot as plt
# plot 4 images as gray scale
plt.subplot(221)
plt.imshow(X_train[0], cmap=plt.get_cmap('gray'))
plt.subplot(222)
plt.imshow(X_train[1], cmap=plt.get_cmap('gray'))
plt.subplot(223)
plt.imshow(X_train[2], cmap=plt.get_cmap('gray'))
plt.subplot(224)
plt.imshow(X_train[3], cmap=plt.get_cmap('gray'))

# show the plot
plt.show()
```



# Code: MNIST on Keras

- Defining our model and model parameters

```
model = keras.Sequential()  
model.add(layers.Dense(20, activation='sigmoid', input_shape=(784,)))  
#Input Layer. The model needs to know what input shape it should expect. For this reason,  
#the first layer in a Sequential model needs to receive information about its input shape  
#Only the first need the shape information, because following layers can do automatic shape  
#inference  
The dense layer is simply a layer where each unit or neuron is connected to each neuron in  
the next layer.  
model.add(layers.Dense(20, activation='sigmoid'))  
  
#In the final layer mention the output classes  
model.add(layers.Dense(10, activation='softmax'))  
  
#Model Summary  
model.summary()
```

# Code: MNIST on Keras

- Understanding the shape of our layers

Model: "sequential"

Layer (type)	Output Shape	Param #
=====		
dense (Dense)	(None, 20)	15700
-----		
dense_1 (Dense)	(None, 20)	420
-----		
dense_2 (Dense)	(None, 10)	210
=====		
Total params: 16,330		
Trainable params: 16,330		
Non-trainable params: 0		
-----		

# Code: MNIST on Keras

Compiling by giving:  
loss function, optimizer  
and validation metric

```
# Compiling model : we define loss function, optimizer and validation metric of our choice  
model.compile(loss='mean_squared_error', metrics=['accuracy'])
```

```
# Fit method: actually running our model by supplying our input and validation data  
model.fit(x_train, y_train, epochs=10)
```

Actual Training or  
running by feeding in  
the data

```
Epoch 1/10  
1875/1875 [=====] - 2s 1ms/step - loss: 0.0416 - accuracy: 0.7382  
Epoch 2/10  
1875/1875 [=====] - 2s 1ms/step - loss: 0.0151 - accuracy: 0.9060  
Epoch 3/10  
1875/1875 [=====] - 2s 1ms/step - loss: 0.0116 - accuracy: 0.9266  
Epoch 4/10  
1875/1875 [=====] - 3s 1ms/step - loss: 0.0101 - accuracy: 0.9359  
Epoch 5/10  
1875/1875 [=====] - 3s 1ms/step - loss: 0.0091 - accuracy: 0.9420  
Epoch 6/10  
1875/1875 [=====] - 3s 1ms/step - loss: 0.0084 - accuracy: 0.9472  
Epoch 7/10  
1875/1875 [=====] - 3s 1ms/step - loss: 0.0079 - accuracy: 0.9499  
Epoch 8/10  
1875/1875 [=====] - 2s 1ms/step - loss: 0.0074 - accuracy: 0.9528  
Epoch 9/10  
1875/1875 [=====] - 2s 1ms/step - loss: 0.0071 - accuracy: 0.9551
```

# Other Advantages of Keras

- Biggest advantage is: Easy and fast
  - Friendliness
  - Modularity
  - Extensibility
- Keras provides easy pipelining of our model.
- Very less and tidy code.
- Pre-existing APIs make our work quite easy.



# Conclusion

- The Deep Learning algorithms are really calculation intensive
- There are many deep learning frameworks
- TensorFlow is one such framework and Keras is a high level API on top of it
- Torch is the next best option.